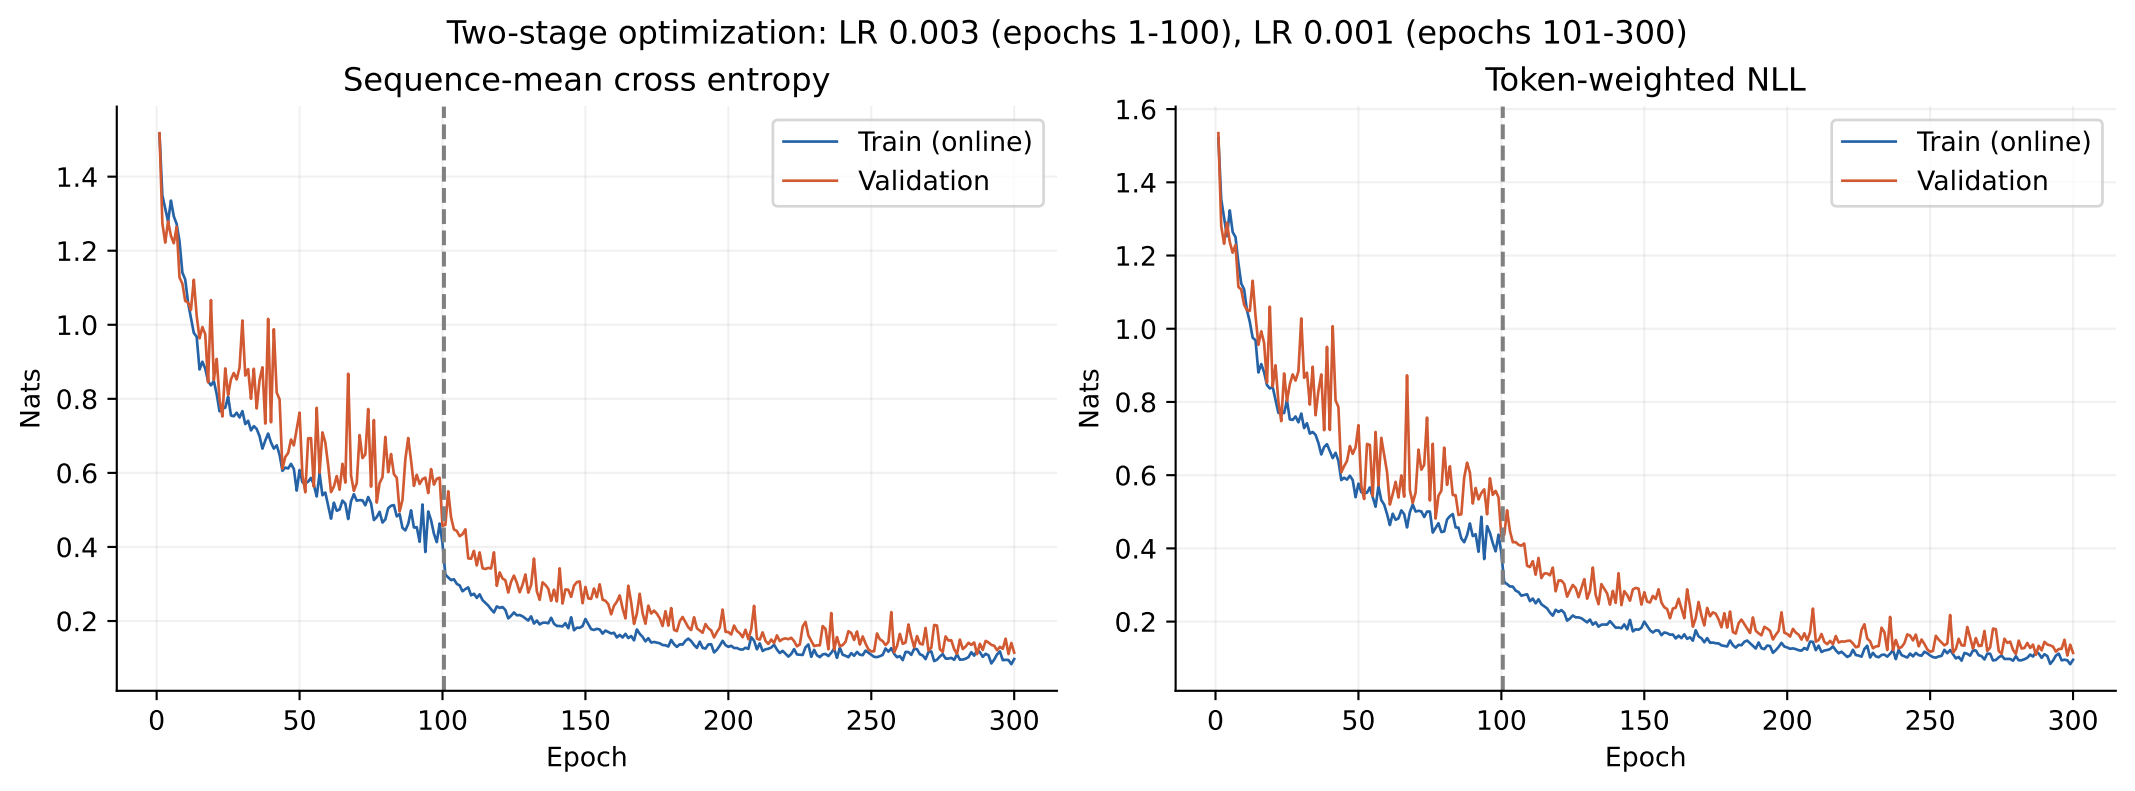
Two-stage optimization: LR 0.003 (epochs 1-100), LR 0.001 (epochs 101-300)
Sequence-mean cross entropy
Token-weighted NLL
Train (online)
Validation
Nats
Epoch
0
50
100
150
200
250
300
0.2
0.4
0.6
0.8
1.0
1.2
1.4
1.6